

Amortized Bayesian Inference with Bayesflow

Paul Bürkner based on materials of Šimon Kucharský

Problem

We want to approximate $p(\theta \mid x)$ or $p(x \mid \theta)$

$$p(\theta \mid x) = \frac{p(\theta) \times p(x \mid \theta)}{p(x)}$$

- Evaluating $p(\theta)$, $p(x \mid \theta)$ and $p(x)$ may not be tractable
- But we are able to sample $(x^{(s)}, \theta^{(s)}) \sim p(x, \theta)$

Solution

Use generative neural networks

- A generative neural network $q_\phi(\theta \mid x)$
 - Normalizing flow, ... (more to come tomorrow)
 - Model the distribution of *parameters*, conditioned on data

Goal:

$$p(\theta \mid x) \approx q_\phi(\theta \mid x)$$

- x conditions the generative network
 - Normalizing flows: Input to a coupling network (usually MLP)
 - Other architectures: Input to a vector field network (usually MLP)

→ Must be of fixed dimensions

- Sample sizes, study design, resolution, ...

→ Summary statistics

1. Handcrafted: Mean, SD, Correlation,...
2. Summary networks

- Takes input of variable size
- Outputs fixed size embedding

→ Condition the inference network on the output of the summary network

$$p(\theta \mid x) \approx q_\phi(\theta \mid h_\psi(x))$$

$$\begin{aligned}\hat{\phi}, \hat{\psi} &= \operatorname{argmin}_{\phi, \psi} \mathbb{E}_{x \sim p(x)} [\mathbb{KL}[p(\theta | x) || q_{\phi}(\theta | h_{\psi}(x))]] \\ &= \operatorname{argmin}_{\phi, \psi} \mathbb{E}_{x \sim p(x)} \left[\mathbb{E}_{\theta \sim p(\theta | x)} \left[\log \frac{p(\theta | x)}{q_{\phi}(\theta | h_{\psi}(x))} \right] \right] \\ &\propto \operatorname{argmin}_{\phi, \psi} - \mathbb{E}_{(x, \theta) \sim p(x, \theta)} [\log q_{\phi}(\theta | h_{\psi}(x))] \\ &\approx \operatorname{argmin}_{\phi, \psi} - \frac{1}{S} \sum_{s=1}^S \log q_{\phi}(\theta^{(s)} | h_{\psi}(x^{(s)}))\end{aligned}$$


→ must be able to generate samples $(x^{(s)}, \theta^{(s)}) \sim p(x, \theta)$

1. Define a generative statistical model
 - $p(x, \theta)$, typically $p(\theta) \times p(x \mid \theta)$
2. Define inference and (optional) summary networks
 - $q_\phi(x \mid h_\psi(x))$
3. Train the networks
 - 3.1 Sample from $(x^{(s)}, \theta^{(s)}) \sim p(x, \theta)$
 - 3.2 Optimize weights ϕ and ψ so that $p(\theta \mid x) \approx q_\phi(\theta \mid h_\psi(x))$


- Once the networks are trained, they can be used for inference
- Most of computational resources are used during training
- Inference is fast (only requires a simple pass through the network)

→ Amortized inference: Pay upfront the cost of inference during training, making subsequent inference effective


Non-Amortized Bayesian Inference



Amortized Bayesian Inference



Breakeven points



MCMC:

- No upfront training
- 5.66 seconds for 10k posterior draws

Amortized Inference:

- 23 seconds upfront training
- 0.07 seconds for 10k posterior draws

Working with bayesflow

- Python library built on `keras3`
 - `TensorFlow`, `JAX`, or `PyTorch` as a backend
- Implementation of common neural architectures to make ABI easier
- Helper functions for simulation, configuration, training, validation, diagnostics, ...
- Website: bayesflow.org
- Github: github.com/bayesflow-org/bayesflow
- Forums: discuss.bayesflow.org
- Install release version: `pip install bayesflow`
- Install dev version: `pip install git+https://github.com/bayesflow-org/bayesflow@dev`

General workflow

1 SIMULATE



Generate data from
any simulation you like

2 AMORTIZE



Part of the  Keras ecosystem
with multi-backend support

3 LEARN



Unleash simulation intelligence
with powerful **generative AI**

- Simulator
 - Simulate data from the statistical model
- Adapter
 - Configure the data
- Summary network (optional)
 - Get summary embeddings
- Inference network
 - Approximate the posterior

- Diagnostics
 - Plots and statistics to evaluate the approximation
- Approximator
 - Hold ingredients necessary for approximation together
 - Adapter, Summary network, Inference network
- Workflow
 - Hold ingredients for the entire pipeline together
 - Simulator, Adapter, Summary network Inference network, Diagnostics

1. Define the statistical model (through simulator)
2. Define the approximator (through neural networks)
3. Train the neural networks
4. Validate the approximator
5. Inference and diagnostics

1. Define the statistical model

Define the simulator

- Define the “statistical” model in terms of a simulator
- `bayesflow` expects “batched” simulations

Listing 1 Python

```
class Simulator(bf.simulators.Simulator):
    def sample(self, batch_size):
        theta = np.random.beta(1, 1, size=batch_size)
        x = np.random.binomial(n=10, p=theta)
        return dict(theta=theta, x=x)

simulator = Simulator()
dataset = simulator.sample(100)
```

Define the simulator

- Define the “statistical” model in terms of a simulator
- `bayesflow` expects “batched” simulations
- `make_simulator`: Convenient interface for “auto batching”

Listing 2 Python

```
def prior():  
    return dict(mu = np.random.normal(0, 1))  
  
def likelihood(mu):  
    return dict(x = np.random.normal(mu, 1, size=20))  
  
simulator = bf.make_simulator([prior, likelihood])  
  
simulator.sample(10)
```

- The shape of data must be the same within each batch

Strategies:

1. Pad vectors to be the same length
2. Vary data shapes between batches

Listing 3 Python

```
def context():
    return dict(n = np.random.randint(10, 101))

def prior():
    return dict(mu = np.random.normal(0, 1))

def likelihood(mu, n):
    observed = np.zeros(0)
    observed[:n] = 1

    x = np.zeros(100)
    x[:n] = np.random.normal(mu, 1, size=n)
    return dict(observed=observed, x=x)

simulator = bf.make_simulator([context, prior, likelihood])
```

Vary data shape between batches

Listing 4 Python

```
def context(batch_size):  
    return dict(n = np.random.randint(10, 101))  
  
def prior():  
    return dict(mu = np.random.normal(0, 1))  
  
def likelihood(mu, n):  
    x = np.random.normal(mu, 1, size=n)  
    return dict(x=x)  
  
simulator = bf.make_simulator(  
    [prior, likelihood], meta_fn=context  
)  
  
simulator.sample(10)
```

Padding vs batched context

- Both approaches work
- Batched context (via `meta_fn`) is a bit cleaner
 - Do not have to fiddle with masking
 - Limited useability (cannot effectively use in offline mode)
- Padding
 - More verbose
 - More general

2. Define the approximator

1. Inference network
 - Generates distributions
2. (Optional) summary network
 - Generates summary embeddings
3. Adapter
 - Reshapes data to be passed into networks

Listing 5 Python

```
inference_network = bf.networks.CouplingFlow()
```

Listing 6 Python

```
inference_network = bf.networks.FlowMatching()
```

Various options available, see the Two Moons Example.

What architecture to pick?

Coupling flow

- Slow training
- Less expressive
- Faster during inference

Flow matching / diffusion models

- Fast training
- More expressive
- Slow inference on CPU

What architecture to pick?

- **Affine coupling** for low-dimensional, uni-modal posteriors
- **Spline coupling** for harder posteriors, including multi-modal
- **Flow matching or diffusion models** for highly complex posteriors
 - or when inference speed is not of concern
 - or if GPUs available

- Extract all relevant information from the variable-sized data
- Output fixed sized embeddings
- “*Sufficient statistics*”
- *Rule of thumb*: output size should be at least $2\times$ the number of parameters

What architecture to pick?

Reflect symmetries in the data

Exchangeable data

- Deep Sets, Set Transformers,...

Listing 7 Python

```
summary_network = bf.networks.DeepSet()
```

Time series, Sequences, ...

- RNNs (e.g., LSTN), Time Series Transformers, ...

Listing 8 Python

```
summary_network = bf.networks.LSTNet()
```

Adapter: Main purpose

- Reshapes data to be passed into networks
- A hub that handles what is passed into what network

Main keywords:

- **"inference_variables"**: What are the variables that the inference network should learn about?
 - $q(\theta \mid x) \rightarrow$ parameters θ
- **"inference_conditions"**: What are the variables that the inference network should be directly conditioned on?
 - $q(\theta \mid x) \rightarrow$ data x
- **"summary_variables"**: What are the variables that are supposed to be passed into the summary network?
 - $q(\theta \mid f(x)) \rightarrow$ data x

Listing 9 Python

```
adapter = (bf.Adapter()  
    .concatenate(["mu", "sigma"], into = "inference_variables")  
    .rename(["n"], "inference_conditions")  
    .concatenate(["x", "observed"], into = "summary_variables")  
)
```

- Transform variables
 - `.standardize`
 - `.sqrt`, `.log`
 - `.constrain`
- Reshape variables
 - `.as_set`, `.as_time_series`
 - `.broadcast`
 - `.one_hot`
- ... and many more operations

- Holds ingredients used for inference

Listing 10 Python

```
approximator = bf.approximators.ContinuousApproximator(  
    inference_network=inference_network,  
    summary_network=summary_network  
    adapter=adapter  
)
```

- Hold all ingredients used for training, inference, and diagnostics

Listing 11 Python

```
workflow = bf.BasicWorkflow(  
    inference_network=inference_network,  
    summary_network=summary_network  
    adapter=adapter,  
    simulator=simulator  
)
```

3. Train the neural networks

Network training

- In principle, training as any other model in `keras`
 - `approximator.fit`
- Define optimizer (e.g., `keras.optimizers.Adam`)
 - (optional) define schedule, early stopping,...
- Define training budget and regime
 - Number of epochs
 - Batch size, Number of batches
 - Simulate data or supply simulator
- Compile and train the model until convergence
- **BasicWorkflow** makes fitting easier, comes with some predefined reasonable settings
 - e.g. `workflow.fit_online`

Online

- Generate data during training
- Slower to train
- Prevents overfitting

Offline

- Train on fixed data
- Faster to train
- Risk of overfitting

From disk

- Offline training
- Large data that does not fit into memory
- Read data on demand during training

4. Validate the approximator

1. Is the neural approximator doing a good job at approximating the true posterior?
 - “Computational faithfulness”
 - Assesses the quality of approximator
2. How much can we expect to learn given data?
 - Parameter recovery, posterior contraction, posterior z-score
 - Assesses how the *statistical* model interacts with data

For general discussion, see Schad et al. (2021).

1. Draw fresh validation data from the simulator
2. Extract the parameter samples from the prior
3. Sample from the posterior

Listing 12 Python

```
dataset = simulator.sample(1_000)
prior = {k: v if k in param_keys for k, v in dataset.items()}
posterior = approximator.sample(500, conditions = dataset)
```

- Testing computational faithfulness of a posterior approximator
- Posterior distribution averaged over prior predictive distribution is the same as the prior distribution

$$p(\theta) = \int \int p(\theta \mid \tilde{y}) \underbrace{p(\tilde{y} \mid \tilde{\theta})p(\tilde{\theta})}_{\text{Prior predictives}} d\tilde{\theta} d\tilde{y}$$

1. Draw N prior predictive datasets $(\theta_i^{\text{sim}}, y_i^{\text{sim}}) \sim p(\theta, y)$
2. For each data set y_i , draw M samples from the approximate posterior:
 $\theta_{ij} \sim q(\theta \mid y_i^{\text{sim}})$
3. Calculate rank statistic $r_i = \sum_{j=1}^M \mathbb{I}(\theta_{ij} < \theta_i^{\text{sim}})$

Simulate

$$\theta^{\text{sim}} \sim p(\theta)$$

$$y^{\text{sim}} \sim p(y \mid \theta^{\text{sim}})$$

By symmetry

$$(\theta^{\text{sim}}, y^{\text{sim}}) \sim p(\theta, y)$$

$$\theta^{\text{sim}} \sim p(\theta \mid y^{\text{sim}})$$

Compute


$$\theta_1, \dots, \theta_M \sim q(\theta \mid y^{\text{sim}})$$

If $q(\theta \mid y^{\text{sim}}) = p(\theta \mid y^{\text{sim}})$, then the rank statistic of θ^{sim} is uniform.

1. Histograms
2. ECDF
3. ECDF difference


Listing 13 Python

```
fig=bf.diagnostics.calibration_histogram(  
    estimates=posterior,  
    targets=prior  
)
```




Listing 14 Python

```
fig=bf.diagnostics.calibration_ecdf(  
    estimates=posterior,  
    targets=prior  
)
```



Listing 15 Python

```
fig=bf.diagnostics.calibration_ecdf(  
    estimates=posterior,  
    targets=prior,  
    difference=True  
)
```



- Necessary but not sufficient condition
 - All faithful approximators should pass the SBC check
 - An unfaithful approximator may pass the SBC check
- Depends on
 - Simulation budget (how many datasets)
 - Approximation precision (how many posterior draws per data set)
 - → degree/ways of miscalibration, rather than ok/not ok

What if SBC check fails?

- Possible causes:
 - Underfitting → increase expressiveness of inference or summary network, increase training budget
 - Overfitting → decrease expressiveness of inference or summary network, increase training budget
 - Coding error → check for errors in the simulator, setting up the approximator
 - Fundamental problem with the statistical model → change the statistical model

- How well does the estimated posterior mean match the true parameter used for simulating the data?

$$z = \frac{\text{mean}(\theta_i) - \theta^{\text{sim}}}{\text{sd}(\theta_i)}$$


- How much uncertainty is removed after updating the prior to posterior?

$$\text{contraction} = 1 - \frac{\text{sd}(\theta_i)}{\text{sd}(\theta^{\text{sim}})}$$

Z-score vs. contraction plot

Listing 16 Python

```
fig=bf.diagnostics.plots.z_score_contraction(  
    estimates=posterior,  
    targets=prior  
)
```




Parameter recovery

- Plot the true parameter values θ^{sim} against a posterior point estimate (mean, median)

Listing 17 Python

```
fig=bf.diagnostics.plots.recovery(  
    estimates=posterior,  
    targets=prior  
)
```



What if I get poor results

- Possible causes:
 - Poor calibration of the approximator → see SBC
 - Poor identification of the parameters
 - Use more informative priors
 - Change the model
 - Add more data (sample size, additional variables)

5. Inference and diagnostics

Obtain posterior samples

Listing 18 Python

```
data = dict(y = ...)
posterior = approximator.sample(1000, conditions=data)

fig=bf.diagnostics.plots.pairs_posterior(estimates=posterior)
```


- Prior misspecification
- Likelihood misspecification

→ Approximator may not be trusted

- Use posterior parameter posterior to simulate data from the simulator
- Compare simulated data to observed data
 - Visual checks - histograms, ecdf plots, scatter plots, ...
 - Comparing summary statistics - means, variances, ...

The Danger of Amortization Gaps

If the real-world test data looks different from the simulated training data, ABI may become inaccurate!



{height

Use summary space for out-of-distribution detection

- Add a base distribution to the summary network, e.g.

Listing 19 Python

```
summary_network = bf.networks.Deepset(base_distribution="normal")
```




Figure 1: Schmitt et al. (2023)

Time for exercise `bayesflow-normal`

Conclusion

In addition to learning $p(\theta \mid x)$:

- Generative modeling tasks (unrelated to Bayesian inference)
- Surrogate likelihood: $p(x \mid \theta)$ (Radev et al., 2023)
- Model comparison: $p(\mathcal{M} \mid x)$ (Elsemüller et al., 2024; Radev et al., 2021)
- Amortized point estimation: $\hat{\theta}$ (Sainsbury-Dale et al., 2024)

Advantages

- Fast inference
- Simulation based – intractable models

Disatvantages

- Need for training
- Simulation based – weaker guarantees

When to use ABI

- Intractable models (models without analytic likelihood density)
- Fast inference needed, e.g.
 - Fitting hundreds of datasets
 - Online estimation (e.g., during an experimental setup)

We do these things not because they are easy, but because we thought they were going to be easy.

— misquote of John F. Kennedy, unknown author

Resources:

- BayesFlow website: documentation, examples
- BayesFlow forum: ask questions about your use cases
- BayesFlow repository: contribute, submit bug reports, feature requests

- Else Müller, L., Schnuerch, M., Bürkner, P.-C., & Radev, S. T. (2024). A deep learning method for comparing bayesian hierarchical models. *Psychological Methods*, Advance online publication. <https://doi.org/10.1037/met0000645>
- Kühmichel, L., others, Bürkner, P.-C., & Radev, S. T. (2026). BayesFlow 2: Multi-Backend Amortized Bayesian Inference in Python. *arXiv Preprint*. <https://doi.org/10.48550/arXiv.2602.07098>
- Radev, S. T., D'Alessandro, M., Mertens, U. K., Voss, A., Köthe, U., & Bürkner, P.-C. (2021). Amortized bayesian model comparison with evidential deep learning. *IEEE Transactions on Neural Networks and Learning Systems*, 34(8), 4903–4917.
- Radev, S. T., Schmitt, M., Pratz, V., Picchini, U., Koethe, U., & Bürkner, P.-C. (2023). JANA: Jointly amortized neural approximation of complex bayesian models. *The 39th Conference on Uncertainty in Artificial Intelligence*. <https://openreview.net/forum?id=dS3wVICQrU0>
- Sainsbury-Dale, M., Zammit-Mangion, A., & Huser, R. (2024). Likelihood-free parameter estimation with neural bayes estimators. *The American Statistician*, 78(1), 1–14.